

```
In [10]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score

from statsmodels.stats.outliers_influence import variance_inflation_factor
from scipy import stats
```

```
In [63]: #Jamboree has helped thousands of students make it to top colleges abroad.
#Be it GMAT, GRE or SAT, their unique problem-solving methods ensure maximum scores with minimum effort.
#They recently launched a feature where students/Learners can come to their website and check their probability
#of getting into the IVY league college. This feature estimates the chances of graduate admission from an India

#Column Profiling:

#Serial No. (Unique row ID)
#GRE Scores (out of 340)
#TOEFL Scores (out of 120)
#University Rating (out of 5)
#Statement of Purpose and Letter of Recommendation Strength (out of 5)
#Undergraduate GPA (out of 10)
#Research Experience (either 0 or 1)
#Chance of Admit (ranging from 0 to 1)
#Problem Statement: Predict the chances of graduate admission based on the given features
```

```
In [4]: df=pd.read_csv("C:\\Users\\prafu\\OneDrive\\Desktop\\Jamboree_Admission.csv")
df
```

Out[4]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65
...	...	...	...	...	...	...	...	...	...
495	496	332	108	5	4.5	4.0	9.02	1	0.87
496	497	337	117	5	5.0	5.0	9.87	1	0.96
497	498	330	120	5	4.5	5.0	9.56	1	0.93
498	499	312	103	4	4.0	5.0	8.43	0	0.73
499	500	327	113	4	4.5	4.5	9.04	0	0.84

500 rows × 9 columns

```
In [5]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Serial No.            500 non-null   int64  
 1   GRE Score              500 non-null   int64  
 2   TOEFL Score            500 non-null   int64  
 3   University Rating      500 non-null   int64  
 4   SOP                    500 non-null   float64 
 5   LOR                    500 non-null   float64 
 6   CGPA                   500 non-null   float64 
 7   Research               500 non-null   int64  
 8   Chance of Admit        500 non-null   float64 
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

In [ ]:

```
In [12]: df.describe()
```





```
Out[12]:
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
count	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000
mean	250.500000	316.472000	107.192000	3.114000	3.374000	3.48400	8.576440	0.560000	0.72174
std	144.481833	11.295148	6.081868	1.143512	0.991004	0.92545	0.604813	0.496884	0.14114
min	1.000000	290.000000	92.000000	1.000000	1.000000	1.00000	6.800000	0.000000	0.34000
25%	125.750000	308.000000	103.000000	2.000000	2.500000	3.00000	8.127500	0.000000	0.63000
50%	250.500000	317.000000	107.000000	3.000000	3.500000	3.50000	8.560000	1.000000	0.72000
75%	375.250000	325.000000	112.000000	4.000000	4.000000	4.00000	9.040000	1.000000	0.82000
max	500.000000	340.000000	120.000000	5.000000	5.000000	5.00000	9.920000	1.000000	0.97000

```
In [11]: cat_cols = ['University Rating', 'SOP', 'LOR ', 'Research']
num_cols = ['GRE Score', 'TOEFL Score', 'CGPA']
target = 'Chance of Admit '
```

```
In [13]: #Univariate Analysis
# check distribution of each numerical variable
rows, cols = 2, 2
fig, axs = plt.subplots(rows,cols, figsize=(12, 8))
index = 0
for row in range(rows):
    for col in range(cols):
        sns.histplot(df[num_cols[index]], kde=True, ax=axs[row,col])
        index += 1
    break


sns.histplot(df[num_cols[-1]], kde=True, ax=axs[1,0])
sns.histplot(df[target], kde=True, ax=axs[1,1])
plt.show()
```



```
In [14]: # check for outliers
rows, cols = 2, 2
fig, axs = plt.subplots(rows, cols, figsize=(12, 7))

index = 0
for col in range(cols):
    sns.boxplot(x=num_cols[index], data=df, ax=axs[0,index])
    index += 1

sns.boxplot(x=num_cols[-1], data=df, ax=axs[1,0])
sns.boxplot(x=target, data=df, ax=axs[1,1])
plt.show()
```



```
In [15]: # There are no outliers present in the dataset.
```


```
In [16]: # check unique values in categorical variables
for col in cat_cols:
    print("Column: {:18} Unique values: {}".format(col, df[col].nunique()))
```

```
Column: University Rating Unique values: 5
Column: SOP Unique values: 9
Column: LOR Unique values: 9
Column: Research Unique values: 2
```

```
In [17]: # countplots for categorical variables
cols, rows = 2, 2
fig, axs = plt.subplots(rows, cols, figsize=(10, 7))

index = 0
for row in range(rows):
    for col in range(cols):
        sns.countplot(x=cat_cols[index], data=df, ax=axs[row, col], alpha=0.8)
        index += 1

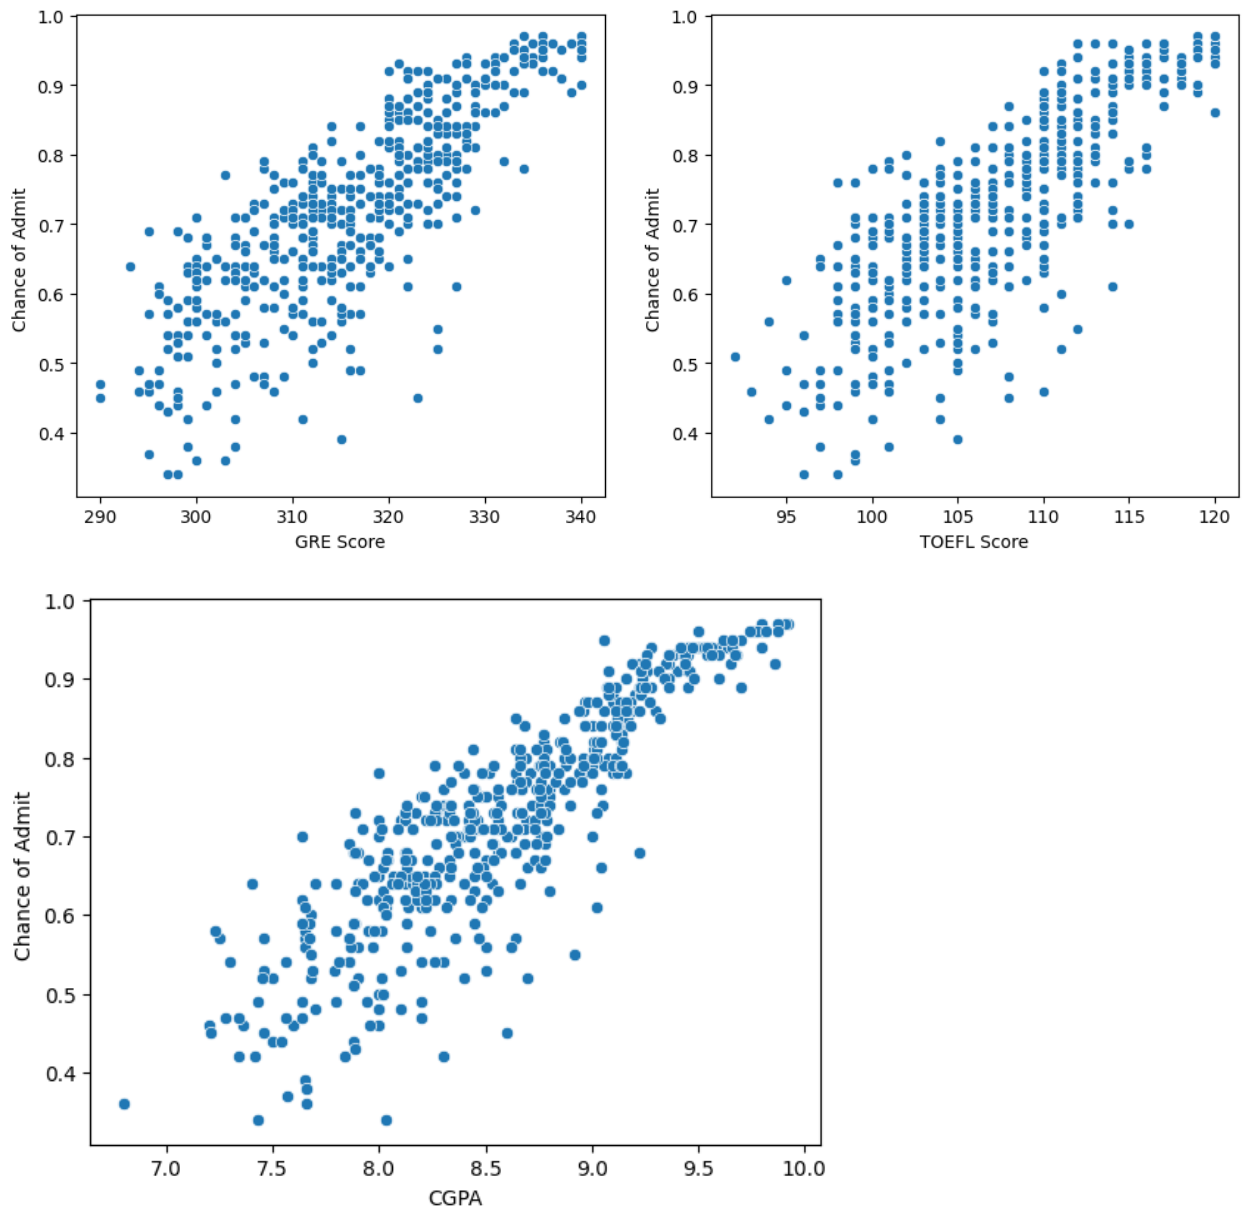
plt.show()
```



```
In [18]: #Bivariate Analysis
```

```
In [19]: # check relation bw continuous variables & target variable
fig, axs = plt.subplots(1, 2, figsize=(12,5))

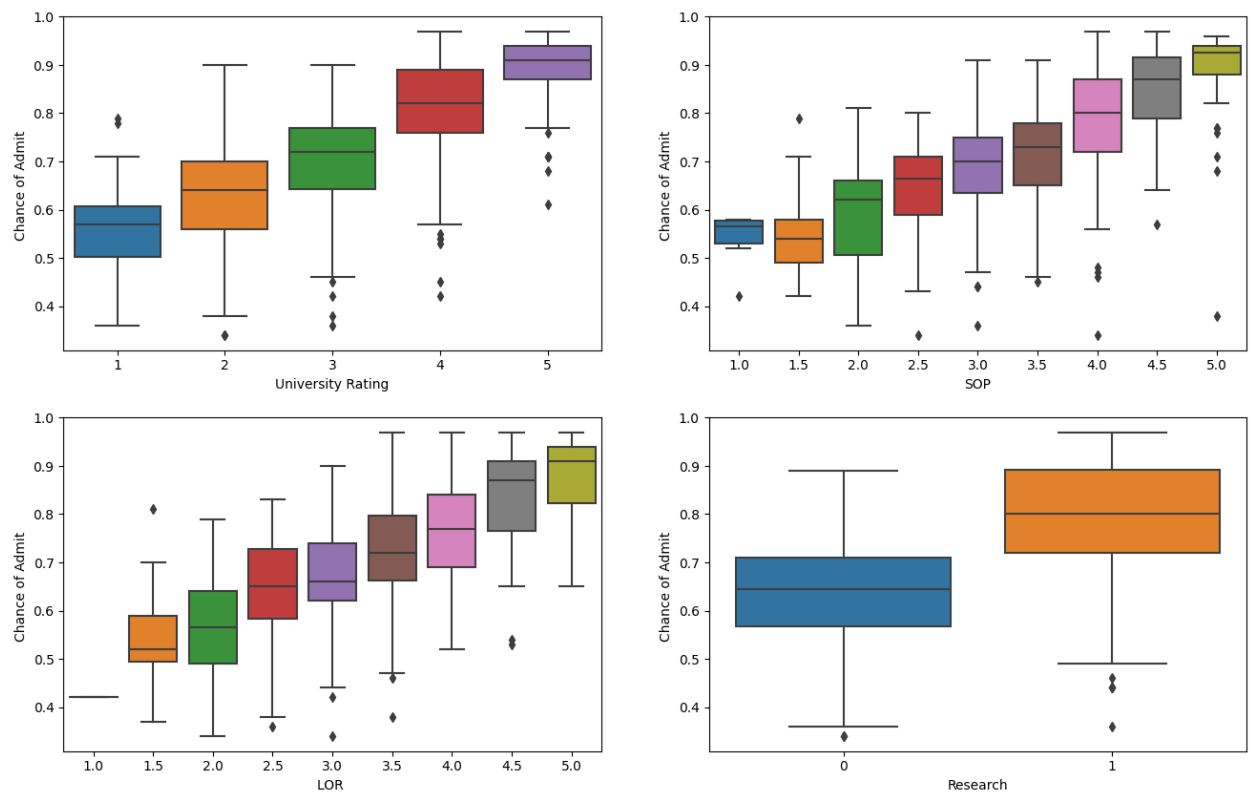
sns.scatterplot(x=num_cols[0], y=target, data=df, ax=axs[0])
sns.scatterplot(x=num_cols[1], y=target, data=df, ax=axs[1])
plt.show()
sns.scatterplot(x=num_cols[2], y=target, data=df)
plt.show()
```



```
In [20]: #there is a linear correlation between the continuous variables and the target variable
```

```
In [21]: rows, cols = 2,2
fig, axs = plt.subplots(rows, cols, figsize=(16,10))


index = 0
for row in range(rows):
    for col in range(cols):
        sns.boxplot(x=cat_cols[index], y=target, data=df, ax=axs[row,col])
        index += 1
```



```
In [22]: #as the rating increases the Chance of Admit also increases.
#Students who have the research experience have more chances of Admin as compared to other students who don't h
```

```
In [23]: #Multivariate Analy
```

```
In [24]: sns.pairplot(df[num_cols])
plt.show()
```



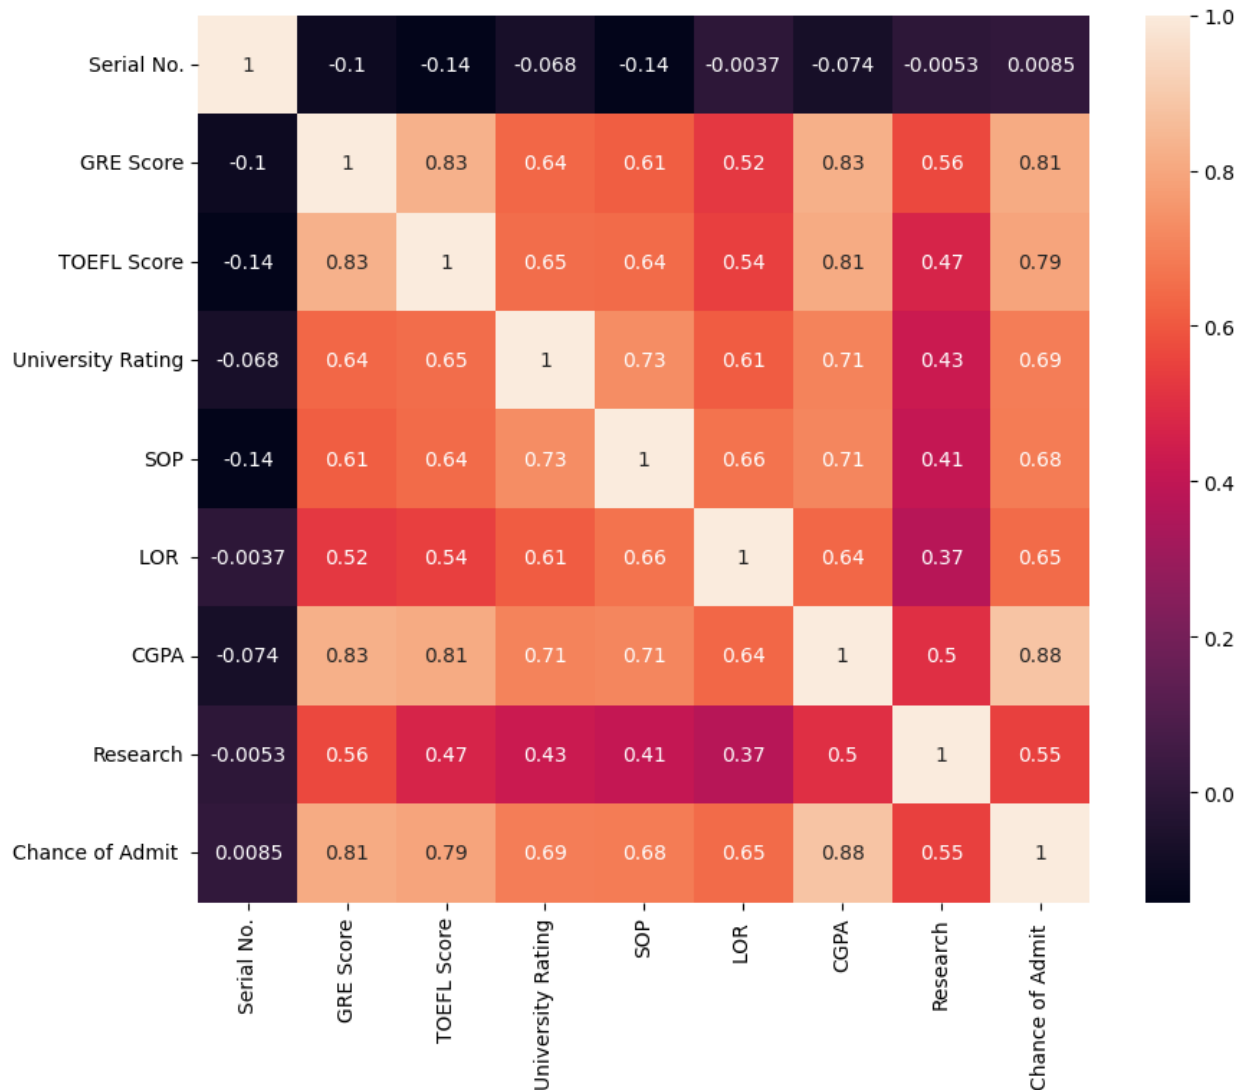
```
In [25]: #Independent continuous variables are also correlated with each other
```

```
In [26]: df.corr()
```

Out[26]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
<b>Serial No.</b>	1.000000	-0.103839	-0.141696	-0.067641	-0.137352	-0.003694	-0.074289	-0.005332	0.008505
<b>GRE Score</b>	-0.103839	1.000000	0.827200	0.635376	0.613498	0.524679	0.825878	0.563398	0.810351
<b>TOEFL Score</b>	-0.141696	0.827200	1.000000	0.649799	0.644410	0.541563	0.810574	0.467012	0.792228
<b>University Rating</b>	-0.067641	0.635376	0.649799	1.000000	0.728024	0.608651	0.705254	0.427047	0.690132
<b>SOP</b>	-0.137352	0.613498	0.644410	0.728024	1.000000	0.663707	0.712154	0.408116	0.684137
<b>LOR</b>	-0.003694	0.524679	0.541563	0.608651	0.663707	1.000000	0.637469	0.372526	0.645365
<b>CGPA</b>	-0.074289	0.825878	0.810574	0.705254	0.712154	0.637469	1.000000	0.501311	0.882413
<b>Research</b>	-0.005332	0.563398	0.467012	0.427047	0.408116	0.372526	0.501311	1.000000	0.545871
<b>Chance of Admit</b>	0.008505	0.810351	0.792228	0.690132	0.684137	0.645365	0.882413	0.545871	1.000000

```
In [27]: plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
In [28]: #Data Preprocessing
```

```
In [29]: # drop Serial NO. column
df = df.drop(columns=['Serial No.'], axis=1)
```

```
In [30]: # check for duplicates
df.duplicated().sum()
```

```
Out[30]: 0
```

```
In [31]: #There are no missing values, outliers and duplicates present in the dataset
```

```
In [32]: #Data preparation for model building
X = df.drop(columns=[target])
y = df[target]
```

```
In [33]: # standardize the dataset
sc = StandardScaler()
X = sc.fit_transform(X)
```

```
In [34]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
```

```
In [35]: print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)

(350, 7) (350,)
(150, 7) (150,)
```



In [36]: *#Model Building*

```
In [38]: def adjusted_r2(r2, p, n):
    """
    n: no of samples
    p: no of predictors
    r2: r2 score
    """
    adj_r2 = 1 - ((1-r2)*(n-1) / (n-p-1))
    return adj_r2

def get_metrics(y_true, y_pred, p=None):
    n = y_true.shape[0]
    mse = np.sum((y_true - y_pred)**2) / n
    rmse = np.sqrt(mse)
    mae = np.mean(np.abs(y_true - y_pred))
    score = r2_score(y_true, y_pred)
    adj_r2 = None
    if p is not None:
        adj_r2 = adjusted_r2(score, p, n)

    res = {
        "mean_absolute_error": round(mae, 2),
        "rmse": round(rmse, 2),
        "r2_score": round(score, 2),
        "adj_r2": round(adj_r2, 2)
    }
    return res
```

```
In [39]: def train_model(X_train, y_train, X_test, y_test, cols, model_name="linear", alpha=1.0):
    model = None
    if model_name == "lasso":
        model = Lasso(alpha=alpha)
    elif model_name == "ridge":
        model = Ridge(alpha=alpha)
    else:
        model = LinearRegression()

    model.fit(X_train, y_train)
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)
    p = X_train.shape[1]
    train_res = get_metrics(y_train, y_pred_train, p)
    test_res = get_metrics(y_test, y_pred_test, p)
    print(f"\n---- {model_name.title()} Regression Model ----\n")
    print(f"Train MAE: {train_res['mean_absolute_error']} Test MAE: {test_res['mean_absolute_error']}")
    print(f"Train RMSE: {train_res['rmse']} Test RMSE: {test_res['rmse']}")
    print(f"Train R2_score: {train_res['r2_score']} Test R2_score: {test_res['r2_score']}")
    print(f"Train Adjusted_R2: {train_res['adj_r2']} Test Adjusted_R2: {test_res['adj_r2']}")
    print(f"Intercept: {model.intercept_}")
    #print(Len(df.columns), Len(model.coef_))
    coef_df = pd.DataFrame({"Column": cols, "Coef": model.coef_})
    print(coef_df)
    print("-"*50)
    return model
```

```
In [40]: train_model(X_train, y_train, X_test, y_test, df.columns[:-1], "linear")
train_model(X_train, y_train, X_test, y_test, df.columns[:-1], "ridge")
train_model(X_train, y_train, X_test, y_test, df.columns[:-1], "lasso", 0.001)
```

---- Linear Regression Model ----

Train MAE: 0.04 Test MAE: 0.04  
 Train RMSE: 0.06 Test RMSE: 0.06  
 Train R2\_score: 0.82 Test R2\_score: 0.82  
 Train Adjusted\_R2: 0.82 Test Adjusted\_R2: 0.81  
 Intercept: 0.724978121476996

	Column	Coef
0	GRE Score	0.018657
1	TOEFL Score	0.023176
2	University Rating	0.011565
3	SOP	-0.000999
4	LOR	0.012497
5	CGPA	0.064671
6	Research	0.013968

---- Ridge Regression Model ----

Train MAE: 0.04 Test MAE: 0.04  
 Train RMSE: 0.06 Test RMSE: 0.06  
 Train R2\_score: 0.82 Test R2\_score: 0.82  
 Train Adjusted\_R2: 0.82 Test Adjusted\_R2: 0.81  
 Intercept: 0.7249823645841696

	Column	Coef
0	GRE Score	0.018902
1	TOEFL Score	0.023252
2	University Rating	0.011594
3	SOP	-0.000798
4	LOR	0.012539
5	CGPA	0.064004
6	Research	0.013990

---- Lasso Regression Model ----

Train MAE: 0.04 Test MAE: 0.04  
 Train RMSE: 0.06 Test RMSE: 0.06  
 Train R2\_score: 0.82 Test R2\_score: 0.82  
 Train Adjusted\_R2: 0.82 Test Adjusted\_R2: 0.81  
 Intercept: 0.7249659139557142

	Column	Coef
0	GRE Score	0.018671
1	TOEFL Score	0.022770
2	University Rating	0.010909
3	SOP	0.000000
4	LOR	0.011752
5	CGPA	0.064483
6	Research	0.013401

Out[40]:

```

└─ Lasso
Lasso(alpha=0.001)
```

```
In [41]: #Since model is not overfitting, Results for Linear, Ridge and Lasso are the same.
#R2_score and Adjusted_r2 are almost the same. Hence there are no unnecessary independent variables in the data
```

```
In [42]: #Linear Regression Model - Assumption Test
#Mutlicollinearity Check
```

```
In [43]: def vif(newdf):
# VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = newdf.columns

# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(newdf.values, i)
                    for i in range(len(newdf.columns))]

return vif_data
```

```
In [44]: res = vif(df.iloc[:, :-1])
res
```

```
Out[44]:
```

	feature	VIF
0	GRE Score	1308.061089
1	TOEFL Score	1215.951898
2	University Rating	20.933361
3	SOP	35.265006
4	LOR	30.911476
5	CGPA	950.817985
6	Research	2.869493

```
In [45]: # drop GRE Score and again calculate the VIF
res = vif(df.iloc[:, 1:-1])
res
```

```
Out[45]:
```

	feature	VIF
0	TOEFL Score	639.741892
1	University Rating	19.884298
2	SOP	33.733613
3	LOR	30.631503
4	CGPA	728.778312
5	Research	2.863301

```
In [46]: # # drop TOEFL Score and again calculate the VIF
res = vif(df.iloc[:, 2:-1])
res
```

```
Out[46]:
```

	feature	VIF
0	University Rating	19.777410
1	SOP	33.625178
2	LOR	30.356252
3	CGPA	25.101796
4	Research	2.842227

```
In [47]: # Now Lets drop the SOP and again calculate VIF
res = vif(df.iloc[:, 2:-1].drop(columns=['SOP']))
res
```

```
Out[47]:
```

	feature	VIF
0	University Rating	15.140770
1	LOR	26.918495
2	CGPA	22.369655
3	Research	2.819171

```
In [48]: # Lets drop the LOR as well
newdf = df.iloc[:, 2:-1].drop(columns=['SOP'])
newdf = newdf.drop(columns=['LOR'], axis=1)
res = vif(newdf)
res
```

```
Out[48]:
```

	feature	VIF
0	University Rating	12.498400
1	CGPA	11.040746
2	Research	2.783179

```
In [49]: # drop the University Rating
newdf = newdf.drop(columns=['University Rating'])
res = vif(newdf)
res
```

Out[49]:

	feature	VIF
0	CGPA	2.455008
1	Research	2.455008

```
In [50]: # now again train the model with these only two features
X = df[['CGPA', 'Research']]
sc = StandardScaler()
X = sc.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
```

```
In [51]: model = train_model(X_train, y_train, X_test, y_test, ['CGPA', 'Research'], "linear")
train_model(X_train, y_train, X_test, y_test, ['CGPA', 'Research'], "ridge")
train_model(X_train, y_train, X_test, y_test, ['CGPA', 'Research'], "lasso", 0.001)
```

---- Linear Regression Model ----

Train MAE: 0.05 Test MAE: 0.05  
 Train RMSE: 0.06 Test RMSE: 0.07  
 Train R2\_score: 0.78 Test R2\_score: 0.81  
 Train Adjusted\_R2: 0.78 Test Adjusted\_R2: 0.81  
 Intercept: 0.7247774222727991

	Column	Coef
0	CGPA	0.112050
1	Research	0.020205

---- Ridge Regression Model ----

Train MAE: 0.05 Test MAE: 0.05  
 Train RMSE: 0.06 Test RMSE: 0.07  
 Train R2\_score: 0.78 Test R2\_score: 0.81  
 Train Adjusted\_R2: 0.78 Test Adjusted\_R2: 0.81  
 Intercept: 0.724783030095277

	Column	Coef
0	CGPA	0.111630
1	Research	0.020362

---- Lasso Regression Model ----

Train MAE: 0.05 Test MAE: 0.05  
 Train RMSE: 0.06 Test RMSE: 0.07  
 Train R2\_score: 0.78 Test R2\_score: 0.81  
 Train Adjusted\_R2: 0.78 Test Adjusted\_R2: 0.81  
 Intercept: 0.7247713356661623

	Column	Coef
0	CGPA	0.111344
1	Research	0.019571

Out[51]:

	Lasso
	Lasso(alpha=0.001)


```
In [52]: #After removing collinear features using VIF and using only two features. R2_score and Adjusted_r2 are
#still the same as before the testing dataset.
```

*#Mean of Residuals*  
*#It is clear from RMSE that Mean of Residuals is almost zero.*

*#Linearity of variables*  
*#It is quite clear from EDA that independent variables are linearly dependent on the target variables.*


*#Normality of Residual*

```
In [53]: y_pred = model.predict(X_test)
residuals = (y_test - y_pred)
sns.histplot(residuals)
plt.show()
```



```
In [54]: #Test for Homoscedasticity
```

```
In [55]: plt.scatter(y_pred, residuals)
plt.show()
```



```
In [56]: #Since the plot is not creating a cone type shape. Hence there is no homoscedasticity present in the data.
```

```
In [57]: #Insights
#1)Multicollinearity is present in the data.
#2)After removing collinear features there are only two variables which are important
#in making predictions for the target variables.
#3)Independent variables are linearly correlated with dependent variables
```

```
In [58]: #Recommendations
#1)CGPA and Research are the only two variables which are important in making the prediction for Chance of Adm
#2)CGPA is the most important varibale in making the prediction for the Chance of Admit.
#3)Following are the final model results on the test data:
RMSE: 0.07
MAE: 0.05
R2_score: 0.81
Adjusted_R2: 0.81
```

In [ ]: